

GPU algorithms for speeding up a 2D Fluid Solver

John Callow

California State University Long Beach

john.j.callow@gmail.com

December 14, 2015

Navier-Stokes (no viscosity or external forces for now)

$$\frac{d\vec{u}}{dt} + \vec{u} \cdot \nabla \vec{u} + \frac{1}{\rho} \nabla p = 0$$

Such that

$$\nabla \cdot \vec{u} = 0$$

Instead of directly solving, split into simpler problems

Advection

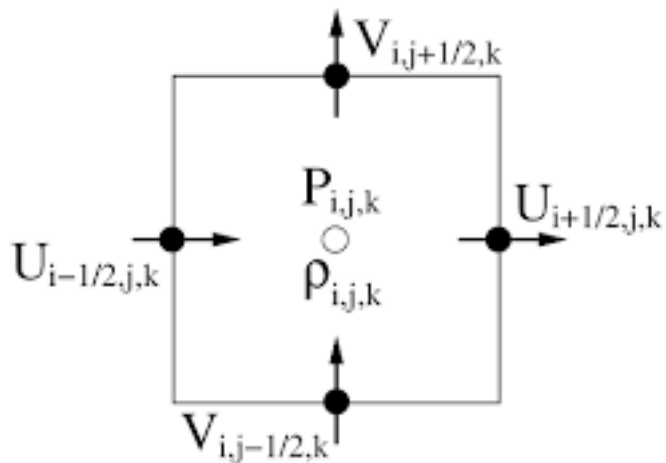
$$\frac{d\vec{u}}{dt} + \vec{u} \cdot \nabla \vec{u} = 0$$

Pressure/Incompressibility

$$\frac{d\vec{u}}{dt} + \frac{1}{\rho} \nabla p = 0$$

Such that

$$\nabla \cdot \vec{u} = 0$$



Finite difference methods unstable. Instead use semi-Lagrangian (Start at grid point, look back in time to see what ends up at this point).

```

1 __global__
2 void advect(GridData* data, float timestep, float delx, GridData *u, GridData *v) {
3     int index = threadIdx.x + blockIdx.x*blockDim.x;
4
5     int x = index % data->width;
6     int y = index/data->width;
7
8     if (x < data->width && y < data->height) {
9
10        float ox = x + data->offset_x;
11        float oy = y + data->offset_y;
12
13        rungeKutta3(data, ox, oy, timestep, delx, u, v);
14
15        data->dst[index] = cubic_interpolation(data, ox, oy);
16    }
17 }

```

Pressure Correction to enforce $\nabla \cdot \vec{u} = 0$

- calculate negative divergence $\vec{b} = -\nabla \cdot \vec{u}_{adv}$
- Set entries of finite difference pressure coefficient matrix A
- Construct MIC(0) - Modified Incomplete Cholesky preconditioner
- Solve $Ap = b$ with MICCG(0) - Conjugate Gradient with MIC(0)
- set $\vec{u}^{n+1} = \vec{u}_{adv} - \Delta t \frac{1}{\rho} \nabla p$

```

1 __global__
2 void buildRHS(GridData *u, GridData *v, float *r, float delx, int
    width, int height) {
3     float scale = 1.0/delx;
4     int id = threadIdx.x + blockIdx.x*blockDim.x;
5
6     int x = id % width;
7     int y = id/width;
8
9     if (x < width && y < height) {
10         r[id] = -scale*(u->get(x+1,y) - u->get(x,y) + v->get(x, y
            +1) - v->get(x,y));
11     }
12 }
13 }

```

Finite Difference applied to pressure update equations

$$u_{i+1/2,j}^{n+1} = u_{i+1/2,j}^{adv} - \Delta t \frac{1}{\rho} \frac{p_{i+1,j} - p_{i,j}}{\Delta x}$$

$$v_{i,j+1/2}^{n+1} = v_{i,j+1/2}^{adv} - \Delta t \frac{1}{\rho} \frac{p_{i,j+1} - p_{i,j}}{\Delta x}$$

Finite Differences to calculate divergence

$$\nabla \cdot \vec{u}_{i,j} = \frac{u_{i+1/2,j} - u_{i-1/2,j}}{\Delta x} + \frac{v_{i,j+1/2} - v_{i,j-1/2}}{\Delta x}$$

Combining the above, taking the divergence, setting to 0, and solving for p

$$\begin{aligned} \frac{\Delta t}{\rho} \left(\frac{4p_{i,j} - p_{i+1,j} - p_{i,j+1} - p_{i-1,j} - p_{i,j-1}}{\Delta x} \right) \\ = - \left(\frac{u_{i+1/2,j} - u_{i-1/2,j}}{\Delta x} + \frac{v_{i,j+1/2} - v_{i,j-1/2}}{\Delta x} \right) \end{aligned}$$

This can be represented as the problem $Ap = b$ with A being a matrix of the above pressure coefficients, p the pressures as a vector, and b the vector of divergence values

Algorithm Preconditioned Conjugate Gradient

```

p = 0
r = b
if infNorm(r) = 0 then
    return p
end if
z = applyPreconditioner(r)
s = z
while iterations < max do
    z = MatrixVectorProduct(A, s)
     $\alpha = \sigma / \text{dotProduct}(z, s)$ 
    p = p +  $\alpha s$ 
    r = r -  $\alpha s$ 
    if infNorm(r)  $\leq$  tol then
        return p
    end if
    z = applyPreconditioner(r)
     $\sigma_{\text{new}} = \text{dotProduct}(z, r)$ 
     $\beta = \sigma_{\text{new}} / \sigma$ 
    s = z +  $\beta s$ 
     $\sigma = \sigma_{\text{new}}$ 
end while
return p

```

```

1 __global__
2 void dotProduct_Device(float *a, float *b, float *result, int ←
    length, int nearest2pow) {
3     extern __shared__ float tmp[];
4     int id = threadIdx.x + blockIdx.x*blockDim.x;
5     int locId = threadIdx.x;
6
7     if (id < length) {
8         tmp[locId] = a[id]*b[id];
9     } else {
10         tmp[locId] = 0;
11     }
12     __syncthreads();
13
14     for (int i = nearest2pow/2; i > 0; i /= 2) {
15         if (locId < i && locId + i < blockDim.x)
16             tmp[locId] += tmp[locId + i];
17         __syncthreads();
18     }
19
20     if (locId == 0)
21         result[blockIdx.x] = tmp[0];
22 }

```

```

1 __host__
2 float dotProduct(float *a, float *b, float* dev_partial, int ←
   length, int blocks, int threads, int nearest2pow) {
3
4   dotProduct_Device<<<blocks, threads, blocks*sizeof(float)>>>(a←
   , b, dev_partial, length, nearest2pow);
5
6   float partial[blocks];
7
8   CUDA_CHECK_RETURN(cudaMemcpy(partial, dev_partial, blocks*←
   sizeof(float), cudaMemcpyDeviceToHost));
9
10  float dotproduct = 0;
11  for (int i = 0; i < blocks; i++) {
12    dotproduct += partial[i];
13  }
14
15  return dotproduct;
16 }

```

```

1 __global__
2 void infinityNorm_Device(float *a, float *result, int length, int ←
    nearest2pow) {
3     extern __shared__ float tmp[];
4
5     int id = threadIdx.x + blockIdx.x*blockDim.x;
6     int locId = threadIdx.x;
7
8     if (id < length)
9         tmp[locId] = abs(a[id]);
10    else
11        tmp[locId] = 0;
12
13    for (int i = nearest2pow/2; i > 0; i /= 2) {
14        if (locId < i && locId + i < blockDim.x) {
15            tmp[locId] = max(tmp[locId + i], tmp[locId]);
16        }
17        __syncthreads();
18    }
19
20    if (locId == 0) {
21        result[blockIdx.x] = tmp[0];
22    }

```

```

1 __host__
2 float infinityNorm(float* a, float *dev_partial, int length, int ←
    blocks, int threads, int nearest2pow) {
3     infinityNorm_Device<<<(blocks, threads, blocks*sizeof(float)←
        >>>(a, dev_partial, length, nearest2pow);
4
5     float partial[blocks];
6
7     CUDA_CHECK_RETURN(cudaMemcpy(partial, dev_partial, blocks*←
        sizeof(float), cudaMemcpyDeviceToHost));
8
9     float infinity_norm = 0;
10    for (int i = 0; i < blocks; i++) {
11        infinity_norm = max(infinity_norm, partial[i]);
12    }
13
14    return infinity_norm;
15 }

```

```

1 __global__
2 void matrixVectorProduct(float *aDiag, float *aPlusX, float *aPlusY, float *dst, float *b, int width, int height) {
3     int id = threadIdx.x + blockIdx.x*blockDim.x;
4     int x = id % width;
5     int y = id/width;
6
7     if (x < width && y < height) {
8         float t = aDiag[id]*b[id];
9
10        if (x > 0)
11            t += aPlusX[id-1]*b[id-1];
12        if (y > 0)
13            t += aPlusY[id - width]*b[id - width];
14        if (x < width - 1)
15            t += aPlusX[id]*b[id + 1];
16        if (y < height - 1)
17            t += aPlusY[id]*b[id + width];
18        dst[id] = t;
19    }
20 }

```

Algorithm Preconditioner Pseudocode $\tau = 0.97$ $\phi = 0.25$ **for** $i = 1$ to *width* **do** **for** $j = 1$ to *height* **do** $px1 = Aplusi[i - 1, j] * precon[i - 1, j]$ $px2 = Aplusi[i, j - 1] * precon[i, j - 1]$ $py1 = Aplusj[i - 1, j] * precon[i - 1, j]$ $py2 = Aplusj[i, j - 1] * precon[i, j - 1]$ $e = A_{diag}[i, j] - px1^2 - py1^2 - \tau(px1py1 + px2py2)$ **if** $e < \sigma * A_{diag}[i, j]$ **then** $e = A_{diag}[i, j]$ **end if** $precon[i, j] = \frac{1}{\sqrt{e}}$ **end for****end for**


```

1 __global__
2 void buildPreconditioner(float *aDiag, float *aPlusX, float *aPlusY, float *preconditioner, int width, int height) {
3     float tau = 0.97;
4     float sigma = 0.25;
5     int id = threadIdx.x;
6     int x, y, matrixId;
7     int limit = max(width, height);
8     for (int i = 0; i < 2*limit; i++) {
9         x = i - id;
10        y = id;
11        matrixId = x + y*width;
12        if (x >= 0 && x < width && y >= 0 && y < height) {
13            float e = aDiag[matrixId];
14            if (x > 0) {
15                float px = aPlusX[matrixId - 1]*preconditioner[matrixId - 1];
16                float py = aPlusY[matrixId - 1]*preconditioner[matrixId - 1];
17                e = e - (px*px + tau*px*py);
18            }

```

```

1
2     if (y > 0) {
3         float px = aPlusX[matrixId - width]*preconditioner[↵
           matrixId-width];
4         float py = aPlusY[matrixId-width]*preconditioner[↵
           matrixId-width];
5         e = e - (py*py + tau*px*py);
6     }
7     if (e < sigma*aDiag[matrixId]) {
8         e = aDiag[matrixId];
9     }
10    preconditioner[matrixId] = 1.0/sqrt(e);
11 }
12 __syncthreads();
13 }
14 }

```

Algorithm Solves $Lq = r$ then $L^T z = q$ to get $L(L^T z) = r$

```

for  $i = 1$  to  $width$  do
  for  $j = 1$  to  $height$  do
     $t = r[i, j] - Aplus[i - 1, j] * precon[i - 1, j] * q[i - 1, j] -$ 
     $Aplus[j[i, j - 1] * precon[i, j - 1] * q[i, j - 1]$ 
     $q[i, j] = t * precon[i, j]$ 
  end for
end for
for  $i = width$  down to  $1$  do
  for  $j = height$  down to  $1$  do
     $t = q[i, j] - Aplus[i, j] * precon[i, j] * z[i + 1, j] -$ 
     $Aplus[j[i, j] * precon[i, j] * z[i, j + 1]$ 
     $z[i, j] = t * precon[i, j]$ 
  end for
end for

```

```

1 __global__
2 void applyPreconditioner(float *aDiag, float *aPlusX, float *aPlusY, float *preconditioner, float *dst, float *a, int width, int height) {
3     int id = threadIdx.x + blockIdx.x*blockDim.x;
4     int x, y, matrixId;
5     int limit = max(width, height);
6
7     for (int i = 0; i < 2*limit; i++) {
8         x = i - id;
9         y = id;
10        matrixId = x + y*width;
11        if (x >= 0 && x < width && y >= 0 && y < height) {
12            float t = a[matrixId];
13            if (x > 0)
14                t -= aPlusX[matrixId - 1]*preconditioner[matrixId - 1]*dst[matrixId - 1];
15
16            if (y > 0)
17                t -= aPlusY[matrixId - width]*preconditioner[matrixId - width]*dst[matrixId - width];
18            dst[matrixId] = t*preconditioner[matrixId];
19        }

```

```

1  __syncthreads();
2  }
3  for (int i = 2*limit-1; i >= 0; i--) {
4      x = i - limit + id;
5      y = limit-1-id;
6      matrixId = x + y*width;
7      if (x >= 0 && x < width && y >= 0 && y < height) {
8          float t = dst[matrixId];
9          if (x < width-1)
10             t -= aPlusX[matrixId]*preconditioner[matrixId]*dst[↵
                matrixId+1];
11          if (y < height-1)
12             t -= aPlusY[matrixId]*preconditioner[matrixId]*dst[↵
                matrixId+width];
13             dst[matrixId] = t*preconditioner[matrixId];
14         }
15     __syncthreads();
16 }
17 }

```

Currently, the gpu apply preconditioner is bad, so I've use host apply for both tests tests. The cpu code tested against is by tunabrain (see sources). His/her code is also based on Robert Bridson's book.

512x512

CPU - 3.1 seconds per frame

GPU - 2.2 seconds per frame

1024x1024

CPU 35 seconds per frame

Gpu 15 seconds per frame



Bridson, Robert. Fluid Simulation for Computer Graphics. Wellesley, MA: A K Peters Ltd. 2008. Print.



Tunabrain, <https://github.com/tunabrain/incremental-fluids>